

Faithfulness as Information Flow: Evaluating and Training Faithful Chain-of-Thought Reasoning

Jinghan Jia^{†,*} Joe Benton[‡] Eric Easley[‡]

[†]Dept. CSE, Michigan State University

[‡]Anthropic

*Anthropic Fellows Program

Abstract

Chain-of-thought (CoT) reasoning is useful for monitoring language models only when the reasoning trace faithfully reflects the computation that produces the final answer. However, models can rely on prompt-to-answer shortcuts that bypass the CoT, making the visible reasoning trace misleading even when it appears plausible. We study CoT faithfulness through a structural information-flow perspective: faithful reasoning should route answer-relevant information through the mediated path from prompt to CoT to answer, rather than through a direct prompt-to-answer shortcut. This perspective yields a task-agnostic framework based on three complementary properties, sufficiency, completeness, and necessity, which we instantiate with entropy-based, masked-KL, and gradient-based diagnostics. We show that these metrics recover externally judged faithfulness differences in hinted reasoning, and identify a low-entropy failure mode of KL-based diagnostics where gradient-based measures remain more stable. Building on this analysis, we introduce update-time interventions for verifier-based on-policy RL, including attention masking, backward-only gradient masking, CoT gradients, and adversarial perturbations of prompt representations. Across hinted arithmetic, reward-hackable code repair, and DAPO-Math models trained without hints but evaluated under wrong-hint injection, our interventions shift behavioral and structural indicators toward stronger CoT mediation. In particular, they make shortcut and reward-hacking behavior more transparent in the CoT and improve task-agnostic faithfulness metrics, while in some settings also reducing wrong-hint susceptibility. Our results suggest that controlling information flow during training is a practical route toward more faithful and monitorable CoT reasoning. Code is available at <https://github.com/safety-research/faithful-cot>.

1 Introduction

Chain-of-thought (CoT) reasoning has become a central interface for improving and monitoring language-model reasoning [1–3]. By externalizing intermediate steps before the final answer, CoT can make model behavior more inspectable. However, this interface is useful for oversight only when the generated reasoning trace faithfully reflects how the model produced its answer. Prior work shows that models can rely on biased cues, misleading hints, or other hidden shortcuts while emitting plausible CoTs that do not reveal the true source of the answer [4–7]. In such cases, CoT gives an illusion of transparency: the trace appears reasonable, but answer-relevant information bypasses it.

Existing CoT-faithfulness evaluations often rely on task-specific probes, such as hint injection, counterfactual editing, trace perturbation, causal-mediation tests, or verbalization-based criteria [4–6, 8, 9]. Recent work further studies faithfulness decay, simulator-based faithfulness training, information-theoretic monitorability, and benchmark-based monitorability evaluation [10–13]. These studies provide important evidence that CoTs can be unfaithful, but many evaluations remain tied

to particular task formats or anticipated failure modes. This raises a more general question: *can we measure and improve CoT faithfulness without designing a new probe for every setting?*

We address this question by viewing CoT faithfulness as an information-flow problem. For a reasoning trajectory $x = (P, C, A)$, with prompt P , chain-of-thought C , and answer A , faithful reasoning should route answer-relevant information through the mediated path $P \rightarrow C \rightarrow A$. Unfaithful reasoning instead allows a direct shortcut $P \rightarrow A$ that bypasses the CoT. Under this view, a faithful CoT should be *sufficient* for predicting the answer, *complete* in capturing answer-relevant prompt information, and *necessary* for producing the answer.

Based on this perspective, we develop task-agnostic diagnostics for CoT faithfulness using entropy, structural masking, and gradients. We further introduce update-time interventions for verifier-based on-policy RL [14–16] that discourage direct prompt-to-answer shortcut learning while preserving the standard rollout and reward pipeline. Across hinted arithmetic, reward-hackable code repair, and DAPO-Math evaluation under wrong-hint injection, our methods shift both behavioral and structural indicators toward stronger CoT mediation. We summarize our **contributions** as follows:

- We formulate CoT faithfulness as a task-agnostic information-flow problem based on sufficiency, completeness, and necessity.
- We propose entropy-based, masked-KL, and gradient-based diagnostics for measuring whether answer-relevant information is routed through the CoT.
- We introduce structural interventions for on-policy RL, including attention masking, backward-only gradient masking, CoT-only gradients, and adversarial prompt-representation perturbations.
- We show that these interventions improve CoT faithfulness across hinted arithmetic, reward-hackable code repair, and held-out wrong-hint evaluation on DAPO-Math, making shortcut and reward-hacking behavior more visible in the CoT.

2 Related Work

Measuring chain-of-thought faithfulness. Chain-of-thought (CoT) reasoning improves multi-step problem solving, but the reasoning trace need not faithfully reflect the computation behind the final answer. Prior work evaluates CoT faithfulness via biased prompts, counterfactual edits, perturbations, truncations, causal-mediation tests, unlearning, and verbalization criteria [4–9], with recent extensions to faithfulness decay, simulatability, instance-level detection, and LLM-judge reliability [10, 17, 18]. These provide behavioral and benchmark-specific evidence of CoT unfaithfulness. In contrast, we formulate faithfulness as a task-agnostic information-flow property: answer-relevant information should be routed through the mediated path from prompt to CoT to answer.

Improving CoT faithfulness and monitorability. Recent work moves from evaluating to improving CoT monitorability. Counterfactual simulation training rewards CoTs that help a simulator predict model behavior under counterfactual inputs [11]. Information-theoretic approaches propose objectives that improve monitor accuracy [12]. Other studies examine whether monitorability emerges under RL with verifiable rewards, whether models can control or hide their CoTs, and when CoT-level optimization helps or harms monitorability [13, 19–21]. Our work is orthogonal along the design axis: prior approaches modify the *reward* signal and typically require auxiliary models in the training loop, a counterfactual simulator [11], an LLM-judge or trained verbalization classifier [22], or a monitor whose accuracy is the optimization target [12]. We instead modify the *policy update*, leaving rollouts, rewards, and advantages unchanged (Figure A1). Our interventions thus operate under a single verifiable reward with no auxiliary simulator, judge, or monitor, and apply directly in settings where such components are not given (e.g., hinted arithmetic, buggy-code repair). The two axes do not conflict, so our methods are in principle composable with reward-level objectives.

Reward hacking and transparent reasoning. Outcome-based optimization can incentivize models to exploit proxy rewards rather than solve the intended task [23–27]. Recent work studies whether CoT monitoring can expose reward hacking or hidden shortcuts in reasoning models [22, 28–30]. Our experiments connect this literature to training-time structure: standard RL can increase reward while leaving shortcut use under-verbalized, whereas our interventions make shortcut behavior more CoT-mediated and transparent.

3 A Task-Agnostic Framework for CoT Faithfulness

A motivating failure and an information-flow view. Chain-of-thought (CoT) is useful for monitoring only if it faithfully reflects how the model arrives at its answer. Figure 1 (left) shows a representative failure: when a misleading hint is added to the prompt, the model changes its final answer to match the hint while the CoT remains nearly unchanged and never acknowledges the hint’s influence. Behavioral probes, hint injection, CoT truncation, or trace edits, can expose such failures, but each is tied to a particular task or shortcut and must be designed per setting. We instead seek a task-agnostic view, framing faithfulness as a question of *information flow*. For a reasoning trace $\mathbf{x} = (P, C, A)$ consisting of prompt, CoT, and answer, faithful reasoning requires that the prompt influence the answer through the mediated path $P \rightarrow C \rightarrow A$ (Figure 1, right). Unfaithful reasoning admits a direct shortcut $P \rightarrow A$ that bypasses the CoT: the trace may look plausible or even correlate with the answer, but it no longer represents the computation that produced it.




Figure 1: **Information flow in reasoning: faithful CoT vs. shortcut solutions.** **Left:** A misleading hint changes the model’s answer, while the CoT does not reveal the hint’s influence, indicating a prompt-to-answer shortcut. **Right:** In faithful reasoning, answer-relevant information should flow through the mediated path $P \rightarrow C \rightarrow A$. In unfaithful reasoning, the model can additionally rely on a direct shortcut $P \rightarrow A$ that bypasses the CoT.

This view yields three complementary requirements for a faithful CoT, capturing respectively the information content of C , the absence of a $P \rightarrow A$ shortcut, and the causal role of C in producing A .

Sufficiency: the CoT determines the answer. A CoT is *sufficient* if conditioning on it reduces uncertainty about the answer:

$$\text{Sufficiency} \iff H(A | C) \text{ is low.} \quad (1)$$

Sufficiency is an informational property of the $C \rightarrow A$ relation: it asks whether the trace *could* in principle determine the answer. A vague, generic, or off-topic CoT leaves $H(A | C)$ high. Sufficiency does not yet imply that the model relies on C — only that the relevant information is present.

Completeness: the CoT captures all answer-relevant prompt information. A CoT is *complete* if, given C , the prompt provides little additional information about the answer:

$$\text{Completeness} \iff I(P; A | C) \approx 0. \quad (2)$$

When this conditional mutual information is near zero, C screens off P from A : any prompt content relevant to the answer has been absorbed into the trace. A completeness failure indicates a residual $P \rightarrow A$ shortcut. The hinted example in Figure 1 is paradigmatic: the hint changes the answer but its influence does not appear in the CoT.

Necessity: the model uses its CoT. A CoT is *necessary* if perturbing it changes the answer:

$$\text{Necessity} \iff A \text{ depends causally on } C. \quad (3)$$

Necessity distinguishes genuine reasoning from post-hoc rationalization: a CoT may appear sufficient and correlate strongly with the answer while the model in fact answers from P directly. Unlike sufficiency and completeness, necessity is interventional and cannot be read off transcripts alone, it requires perturbing C and observing the effect on A .

The three properties are complementary. The properties are logically independent: a CoT can be sufficient without being necessary (correlated with the answer while not causing it), can satisfy completeness degenerately when outputs are uninformative, and can be necessary while still omitting prompt information that matters. Faithful CoT therefore requires all three.




Figure 2: **Task-agnostic faithfulness metrics.** **Left:** attention masks isolate full, CoT-mediated, and prompt-only answer distributions. **Right:** gradient-based metrics compare answer dependence on prompt versus CoT tokens; greater CoT-gradient concentration indicates stronger CoT-mediated reasoning.

4 Measures: Operationalizing Faithfulness

We instantiate the properties of Section 3 with three families of task-agnostic metrics, summarized in Figure 2 and Table 1: an entropy-based metric for sufficiency, KL-based metrics from structural attention masking for completeness and necessity, and gradient-based metrics from local answer sensitivity for the same two properties.

Sufficiency. We measure sufficiency using the answer entropy under an external reference model q :

$$\text{Suff} = H_q(A | C). \quad (4)$$

Using q rather than the CoT-generating model avoids circularity and asks whether C is informative to an external observer; lower entropy indicates a more sufficient CoT.

KL-based metrics. Let $p(A | P, C)$ denote the answer distribution under standard causal attention. Following Figure 2, structural masks either block the direct $A \leftarrow P$ path (yielding $p(A | C)$) or block answer access to C (yielding $p(A | P)$). KL divergences between the masked and unmasked distributions operationalize completeness and necessity:

$$\text{KL-DE} = D_{\text{KL}}(p(A | P, C) \| p(A | C)), \quad \text{KL-Nec} = D_{\text{KL}}(p(A | P, C) \| p(A | P)). \quad (5)$$

Reference-model variants $\text{KL-DE}_{\text{ref}}$ and $\text{KL-Nec}_{\text{ref}}$ replace p with q (Table 1).

Gradient-based metrics. We complement these finite-difference metrics with local sensitivity. Let $\mathcal{L}_A$ be the answer loss under full causal attention and let $g_t = \|\partial \mathcal{L}_A / \partial e_t\|_2$. Aggregating gradient mass over prompt and CoT positions gives

$$\text{Grad-DE} = \frac{\sum_{t \in P} g_t}{\sum_{t \in P \cup C \cup A} g_t}, \quad \text{Grad-Nec} = \frac{\sum_{t \in C} g_t}{\sum_{t \in P \cup C \cup A} g_t}. \quad (6)$$

A higher Grad-DE indicates greater shortcut reliance; a higher Grad-Nec indicates stronger dependence on the CoT. Since exact interventional necessity is intractable during training, g_t serves as a first-order linearization of the answer’s sensitivity to perturbing token t .

Validation against an external criterion.

To validate the metrics, we compare Qwen3-8B and DeepSeek-R1-Distill-14B on hinted GPQA, restricting to the subset of examples where adding a hint changes the model’s answer. We define an external faithfulness criterion via *verbalized hint-following*: whether the CoT explicitly acknowledges the hint’s influence on the final answer, as judged by four frontier LLMs (GPT-5.4, GPT-5.2, Claude Opus 4.6, Claude Sonnet 4.5). This yields a clear ordering: 89.4% verbalized faithful for Qwen3-8B versus 54.3% for DeepSeek-R1-Distill-14B. We then ask whether our metrics recover this ordering without access to verbalization labels. As shown in Figure 3, the entropy- and gradient-based metrics

Table 1: Task-agnostic faithfulness metrics. p denotes the CoT-generating model.

Property	Metric	Definition	Better
Suff.	Entropy	$H_q(A C)$	↓
Comp.	KL-DE	$D_{\text{KL}}(p(A P, C) \ p(A C))$	↓
	KL-DE _{ref}	$D_{\text{KL}}(q(A P, C) \ q(A C))$	
	Grad-DE	$\frac{\sum_{t \in P} g_t}{\sum_t g_t}$	
Nec.	KL-Nec	$D_{\text{KL}}(p(A P, C) \ p(A P))$	↑
	KL-Nec _{ref}	$D_{\text{KL}}(q(A P, C) \ q(A P))$	
	Grad-Nec	$\frac{\sum_{t \in C} g_t}{\sum_t g_t}$	




Figure 3: **Validation of the proposed task-agnostic faithfulness metrics.** DeepSeek-R1-Distill-14B vs. Qwen3-8B on answer-changing hinted GPQA examples, where an external verbalization criterion indicates Qwen3-8B is more faithful. Bars show mean metric values across validation examples; error bars denote standard error. All metrics significantly distinguish the two models under a two-sided Mann–Whitney U test, but only entropy- and gradient-based metrics recover the predicted direction; the direction of KL-Nec is inverted. Stars denote significance: *** $p < 0.001$; ** $p < 0.01$; * $p < 0.05$. ✓ correct direction (metric ranks the more-faithful model as expected); × inverted direction.

recover the ordering in the direction predicted by the framework: the more faithful model has lower $H_q(A | C)$, lower Grad-DE, and higher Grad-Nec, all significant under a two-sided Mann–Whitney U test on sample-level values. KL-Nec also separates the two models significantly, but in the *wrong direction*: the more faithful model exhibits the smaller CoT-removal shift, contrary to what a clean necessity measure should predict. This suggests KL-Nec carries a bias that is not removed by simple distributional comparison, and is therefore unreliable as a directional faithfulness measure on its own.

KL-DE is confounded in low-entropy regimes. Another issue arises with KL-DE when answer distributions become near-deterministic, as is common once a model has converged on a confident behavior. On a checkpoint trained for buggy-code repair in Sec 6.3, KL-DE depends strongly and non-monotonically on the full-context answer entropy (Figure A2 in Appendix C.2), suggesting that KL is confounded by entropy itself rather than cleanly tracking shortcut reliance. Grad-DE varies much more smoothly over the same range. Despite this limitation, KL-DE remains a useful behavioral signal for detecting faithfulness. Because both metrics only require querying the model’s output distribution under different attention masks, they remain valuable as black-box proxies when model weights are unavailable. We therefore use them as complementary diagnostics and treat gradient-based metrics as the primary faithfulness measures in the remainder of the paper.

5 Structural Interventions for Faithful CoT Training

The metrics above diagnose whether answer-relevant information flows through the CoT; we next ask whether the same information-flow view can guide training. We study structural interventions for GRPO-based on-policy RL. As shown in Figure A1, all methods leave rollout generation, rewards, and advantage normalization unchanged, intervening during the policy update to encourage CoT-mediated rather than direct prompt-to-answer learning, we proposed following intervention methods.

Update mask. Our first intervention blocks direct $A \rightarrow P$ attention during the actor’s update-time forward pass. Let z_{ij} denote the pre-softmax attention logit from query position i to key position j ,

$$z_{ij} \leftarrow -\infty \quad \text{for all } i \in A, j \in P, \quad (7)$$

be set during training, so answer tokens cannot directly attend to prompt tokens when computing the update loss. This removes the most direct prompt-to-answer shortcut during update-time recomputation and encourages the model to rely more on CoT-mediated information.

Gradient mask. A limitation of the update mask is that it changes the actor’s forward computation during the update. Our second intervention avoids this mismatch by leaving the forward pass unchanged and blocking only the backward signal through $A \rightarrow P$ edges. Let $G_{ij} = \mathbb{1}[i \in A \wedge j \in P]$. We replace the relevant logits with

$$\tilde{z}_{ij} = \text{sg}(G_{ij} z_{ij}) + (1 - G_{ij}) z_{ij}, \quad (8)$$

where $\text{sg}(\cdot)$ is the stop-gradient operator. This preserves the standard forward computation while preventing the policy gradient from directly reinforcing prompt-to-answer attention paths.

CoT Gradient. The previous two interventions target specific attention edges. Our third intervention instead targets which token positions are allowed to shape the parameter update. Let $\mathbf{m}$ be the binary mask over CoT positions. For a linear layer $\mathbf{Y} = \mathbf{X}\mathbf{W} + \mathbf{b}$, we use

$$\mathbf{Y} = (\mathbf{X} \odot \mathbf{m})\mathbf{W} + (\mathbf{X} \odot (\mathbf{1} - \mathbf{m}))\text{sg}(\mathbf{W}) + \mathbf{b}, \quad (9)$$

which is numerically identical to the standard forward pass, but ensures that only CoT positions contribute directly to parameter gradients. Intuitively, this concentrates the learning signal on improving CoT representations and how they support the final answer.

FACT. Finally, we propose **FACT** (Faithfulness-Adversarial CoT Training), which perturbs prompt hidden states and trains the model to remain robust to that perturbation during the update. At a chosen layer ℓ , we compute an FGSM-style perturbation

$$\delta^* = \varepsilon \cdot \text{sign}(\nabla_{\delta} \mathcal{L}_{\text{inner}}) \Big|_{\delta=0}, \quad (10)$$

where $\mathcal{L}_{\text{inner}}$ is the teacher-forced policy update loss computed on the perturbed hidden states. We inject δ^* into the prompt representations and then recompute the update loss using the perturbed hidden states.

Table 2: Summary of structural interventions. All methods preserve rollout generation, rewards, and advantage normalization, and differ only in how they modify the GRPO policy update.

Method	Intervention point	Flash attn	Compute overhead	Memory overhead
Standard update	—	✓	—	—
Update mask	Policy fwd: $\alpha_{ij} = 0$ for $A \rightarrow P$	✗	$O(s^2 n_h L)$	$O(s^2 n_h L)$
Gradient mask	Policy bwd: $\nabla_{\alpha, \mathbf{v}} = 0$ for $A \rightarrow P$	✗	$O(s^2 n_h L)$	$O(s^2 n_h L)$
CoT Gradient	Policy bwd: $\nabla_{\mathbf{W}}$ only from C	✓	$O(s d^2 L)$	$O(s d L)$
FACT	Inner adversarial step on layers $[\ell, L]$	✓	$O((L - \ell) s d^2)$	$O(p d)$

6 Experiments

6.1 Experimental Setup

Training tasks and datasets. We train models with GRPO-style REINFORCE [14] on three task families, each targeting a distinct faithfulness failure mode. **(1) Hinted arithmetic.** We use `arithmetic_mul_div_multiple` from the DeepMind Mathematics Dataset [31], augmented with a hint of the form “A Stanford professor guessed {hint}” that is correct with probability 0.75, following Turpin et al. [4]. Reward is +1 for numerically correct answers. The targeted failure mode is a *completeness violation*: answering via the hint without verbalizing it in the CoT. **(2) Buggy code repair.** Derived from Code Contests [32], the model receives buggy code together with visible test assertions and must return fixed code; reward is tied to passing the visible tests. The targeted failure mode is *silent reward hacking*: emitting a hardcoded lookup table that passes the visible tests while the CoT describes genuine bug analysis. **(3) DAPO-math.** We use DAPO-Math-17K [16] for general mathematical reasoning, with no hint or hacking opportunity at training time. We use this setting to probe generalization: at test time we inject wrong-valued hints into a held-out set and measure whether the model’s answer changes.

Models and Methods We train GEMMA3-1B-IT [33] on hinted arithmetic, GEMMA3-4B-IT and GEMMA3-12B-IT on buggy code repair (to assess scaling), and QWEN2.5-7B-INSTRUCT [34] on DAPO-math [16]. We compare standard GRPO (**Vanilla**) against our four structural interventions: **Update mask**, **Gradient mask**, **CoT-only gradients**, and **FACT** (Section 5). Details appear in Appendix B.

Evaluation metrics. We evaluate along two axes: *faithfulness* and *task performance*. For faithfulness, we report the gradient-based metrics Grad-DE and Grad-Nec as primary measures, together with sufficiency $H_q(A | C)$ (Section 4); KL-based metrics, which Section 4 shows are less reliable in the low-entropy regime of well-trained models, are reported in Appendix C.2. We also report task-specific behavioral metrics. For hinted math: the *hint mention rate* (fraction of CoTs containing the hint value) and the *wrong-hint following rate* (fraction of incorrect-hint cases where the answer matches the hint). For code repair: the *lookup-table rate* (fraction of solutions that are hardcoded lookup tables) and the *lookup-intent verbalization rate* (fraction of lookup-table solutions whose <think> block also contains lookup code), both via rule-based pattern matching; complementary LLM-judge results [35] are in Appendix C.3. For DAPO-math: the *hint-influence rate* (fraction of held-out cases where the answer changes under an injected hint).

6.2 Hinted Math: Interventions Make CoT Faithful




Figure 4: Training dynamics of vanilla RL and faithfulness-oriented interventions (CoT Gradient, Gradient Mask, Update Mask, and FACT) in transparency, robustness to wrong hints, and task performance on hinted arithmetic. The four panels report implicit hint mention rate, wrong-hint following rate, accuracy on wrong-hint-conditioned examples, and overall accuracy vs. training steps. Higher implicit hint mention, lower wrong-hint following, higher wrong-hint-conditioned accuracy, and higher overall accuracy indicate better behavior.

Interventions improve the faithfulness of CoT. The hinted arithmetic setting cleanly separates task performance from faithful reasoning: because 75% of training hints are correct, a model that simply learns to follow the hint can reach $\sim 75\%$ accuracy without solving arithmetic robustly. Whether the model exploits this shortcut, and whether it verbalizes doing so, is what we ask of vanilla RL and the four interventions.

Behavioral dynamics (Fig. 4). The fourth panel confirms that vanilla RL and most interventions plateau near the $\sim 75\%$ ceiling expected from hint-following, so overall accuracy alone is uninformative. The mechanism is visible in the second and first panels. Vanilla RL’s wrong-hint following rate climbs sharply across training, indicating that the model exploits the statistical reliability of the hint, while its hint mention rate remains low and unstable: the shortcut is learned but not verbalized, a completeness failure.

Causal information flow (Fig. 5).

The four interventions all substantially raise the hint mention rate, exposing the shortcut when it occurs. The third panel further shows that CoT Gradient attains substantially higher accuracy on wrong-hint examples than vanilla RL or the other interventions, suggesting it not only makes hint use visible but partly prevents reliance on it.

To distinguish surface verbalization from causal influence, we perturb the hint either in the prompt or in the CoT and measure the KL shift in the answer distribution, summarized as $(|KL_{CoT}| - |KL_{prompt}|) / (|KL_{CoT}| + |KL_{prompt}|)$; positive values indicate CoT-dominant routing, negative values prompt-dominant shortcut reliance.

Under sign-flip, $2\times$ scaling, and random hint changes, vanilla RL trends increasingly negative as training proceeds, consistent with the prompt-dominant shortcut suggested by the behavioral evidence. CoT Gradient, Gradient Mask, and Update Mask remain near or above zero across perturbation types; FACT improves over vanilla but less uniformly. The interventions therefore do not merely add surface verifications of the hint, they shift answer determination onto the mediated path through the CoT.




Figure 5: Causal-effect and necessity analysis on hinted arithmetic. We report the normalized KL difference between CoT-hint and prompt-hint sensitivity across training steps under sign-flip, scale- $2\times$, and random hint-value changes. Positive values indicate CoT-dominant behavior, while negative values indicate prompt-dominant shortcut reliance; higher is better.




Figure 6: Faithfulness metrics on hinted arithmetic. We report sufficiency, completeness, and necessity using $H(A | C)$, Grad-DE, and Grad-Nec vs. training steps. Lower $H(A | C)$, lower Grad-DE, and higher Grad-Nec indicate better faithfulness.

Task-agnostic metrics (Fig. 6). The structural metrics from Section 4 agree with the behavioral and causal results. Compared with vanilla RL, the interventions yield lower $H(A | C)$ (more sufficient CoTs), lower Grad-DE (less direct prompt-to-answer dependence), and higher Grad-Nec (greater answer dependence on the CoT). CoT Gradient gives the clearest improvement on completeness and necessity, while the other interventions provide partial gains. Quantitative transparency and faithfulness analysis in Appendix C.5.

6.3 Code Repair: Reward Hacking Becomes Transparent




Figure 7: Behavioral dynamics of vanilla RL and faithfulness-oriented interventions on buggy-code fixing. The four panels report visible/hidden test case pass rate, lookup-table hack rate, and CoT hack verbalization rate vs. training steps. Visible and hidden test pass rates measure reward optimization and generalization, respectively; lookup-table hack rate measures shortcut exploitation of visible tests; and CoT hack verbalization rate measures whether the shortcut strategy is exposed in the CoT. We remark that higher CoT hack verbalization indicates better transparency of the model’s final code behavior.

Interventions expose reward hacking in code repair. Behavioral dynamics (Fig. 7). With reward tied only to visible test assertions, the model can either fix the code or hardcode a lookup table that passes those tests while failing on hidden ones. Across all methods the visible test pass rate rapidly approaches 1, while hidden test accuracy plateaus around 0.32 and the lookup-table hack rate rises in step, confirming that the visible-test gain comes from shortcut learning rather than robust repair. The diagnostic question is whether the CoT exposes this. Vanilla RL almost never verbalizes the lookup-table strategy despite producing such solutions frequently, and Gradient Mask shows the same failure. CoT Gradient, Update Mask, and FACT substantially raise the CoT hack verbalization rate, with CoT Gradient the most stable; verbalization climbs sharply as the model begins to exploit the lookup-table strategy. The interventions do not eliminate reward hacking here, but they make it monitorable from the trace.




Figure 8: Faithfulness metrics on buggy-code fixing. The three panels report sufficiency, completeness, and necessity using $H(A | C)$, Grad-DE, and Grad-Nec across training steps. Better faithfulness corresponds to more sufficient CoT information, weaker direct prompt-to-answer reliance, and stronger answer dependence on the CoT.

Task-agnostic metrics (Fig. 8). The structural metrics agree with the behavioral picture. CoT Gradient achieves the lowest $H(A | C)$ (most sufficient CoTs); CoT Gradient, Update Mask, and FACT reduce Grad-DE relative to vanilla RL and Gradient Mask (weaker residual prompt-to-answer dependence); Update Mask and FACT achieve the largest Grad-Nec, while CoT Gradient remains more stable than vanilla RL. The methods that increase hack verbalization are the same methods that shift answer-relevant information into the CoT and reduce direct shortcut reliance. We observe consistent trends on GEMMA3-12B-IT (Figure A5 and Figure A6 in Appendix C.4).

6.4 DAPO-Math: Faithfulness Generalizes Beyond the Training Distribution

Faithfulness transfers to unseen hint manipulation. The previous two settings show that interventions improve faithfulness on the very tasks they were trained on. A stronger test is whether these gains transfer to manipulations the model never encountered during training. We train QWEN2.5-7B-INSTRUCT on DAPO-Math-17K [16] with vanilla RL and CoT Gradient *without any hint signal*




Figure 9: Behavioral evaluation of vanilla RL and CoT Gradient on DAPO-Math. Models are trained on the DAPO-Math training set and evaluated on 1,000 validation examples. The four panels report no-hint accuracy, wrong-hint follow rate, explicit verbalization rate among hint-following responses, and implicit verbalization rate among hint-following responses vs. training steps. Higher no-hint accuracy indicates stronger task performance, lower wrong-hint follow rate indicates better robustness to incorrect hints, and higher verbalization rates indicate that hint reliance is more explicitly reflected in the CoT. *during training*, then inject wrong-valued hints at evaluation time. Update Mask and Gradient Mask are infeasible at this scale due to their $O(s^2)$ cost.

Behavioral dynamics (Fig. 9). CoT Gradient matches vanilla RL on no-hint accuracy throughout training (panel A), confirming that the intervention preserves capability on a normal mathematical reasoning task at 7B scale and does not impose a tax on standard task performance. Both methods follow wrong hints at comparable rates throughout training (panel B), indicating that neither training procedure makes the model robust to test-time hint injection per se. The diagnostic difference appears *when the model follows*: vanilla RL’s explicit and implicit verbalization rates on follower examples both degrade over training (panels C, D), so hint-following becomes increasingly silent and harder to detect from the trace alone. CoT Gradient shows the opposite trend, with implicit verbalization saturating around 0.9 and explicit verbalization stable through training. The intervention does not prevent the model from being influenced by hints it never saw during training, but it keeps that influence visible in the CoT, exactly the property that makes CoT monitoring viable as a safety signal.

Task-agnostic metrics (Fig. 10). The three structural axes agree with the behavioral evidence. CoT Gradient yields substantially lower $H(A | C)$ than vanilla RL (more sufficient CoTs), Grad-DE that declines steadily over training (weaker direct prompt-to-answer dependence), and Grad-Nec that rises in step (greater dependence of the answer on the CoT). The same separation holds when restricted to the hint-following subset, indicating that the gap reflects genuine structural change in how the model routes information rather than an averaging artifact driven by non-follower examples. Together with the behavioral results, this shows that the structural improvements from Sections 6.2 and 6.3 transfer to a held-out, unanticipated form of hint manipulation, suggesting the interventions induce faithfulness as a generalizable property rather than memorization of a training-time shortcut.




Figure 10: Faithfulness metrics of vanilla RL and CoT Gradient on DAPO-Math. The three panels report sufficiency, completeness, and necessity using $H(A | C)$, Grad-DE, and Grad-Nec across training steps, evaluated on 1,000 validation examples. Lower $H(A | C)$, lower Grad-DE, and higher Grad-Nec indicate better faithfulness.

7 Conclusion

We studied chain-of-thought faithfulness as an information-flow problem, where answer-relevant information should pass through the CoT rather than bypass it through direct prompt-to-answer shortcuts. We proposed task-agnostic diagnostics based on sufficiency, completeness, and necessity, and introduced structural interventions for on-policy RL that encourage CoT-mediated answer formation. Across hinted arithmetic, reward-hackable code repair, and DAPO-Math, these interventions improved both structural faithfulness metrics and behavioral transparency, making shortcut and reward-hacking behavior more visible in the CoT. Our results indicate that shaping information flow during training is a promising route toward more faithful and monitorable reasoning models. Broader impacts and limitations are discussed in Appendices E and D, respectively.

Bibliography

- [1] Maxwell Nye, Anders Johan Andreassen, Guy Gur-Ari, Henryk Michalewski, Jacob Austin, David Bieber, David Dohan, Aitor Lewkowycz, Maarten Bosma, David Luan, et al. Show your work: Scratchpads for intermediate computation with language models. 2021.
- [2] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.
- [3] Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *Advances in neural information processing systems*, 35:22199–22213, 2022.
- [4] Miles Turpin, Julian Michael, Ethan Perez, and Samuel Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. *Advances in Neural Information Processing Systems*, 36:74952–74965, 2023.
- [5] Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023.
- [6] Debjit Paul, Robert West, Antoine Bosselut, and Boi Faltings. Making reasoning matter: Measuring and improving faithfulness of chain-of-thought reasoning. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024.
- [7] James Chua and Owain Evans. Are deepseek r1 and other reasoning models more faithful? *arXiv preprint arXiv:2501.08156*, 2025.
- [8] Iván Arcuschin, Jett Janiak, Robert Krzyzanowski, Senthooran Rajamanoharan, Neel Nanda, and Arthur Conmy. Chain-of-thought reasoning in the wild is not always faithful. *arXiv preprint arXiv:2503.08679*, 2025.
- [9] Martin Tutek, Fateme Hashemi Chaleshtori, Ana Marasovic, and Yonatan Belinkov. Measuring chain of thought faithfulness by unlearning reasoning steps. In Christos Christodoulopoulos, Tanmoy Chakraborty, Carolyn Rose, and Violet Peng, editors, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 9935–9960, Suzhou, China, November 2025. Association for Computational Linguistics. ISBN 979-8-89176-332-6. doi: 10.18653/v1/2025.emnlp-main.504. URL <https://aclanthology.org/2025.emnlp-main.504/>.
- [10] Donald Ye, Max Loffgren, Om Kotadia, and Linus Wong. Mechanistic evidence for faithfulness decay in chain-of-thought reasoning. *arXiv preprint arXiv:2602.11201*, 2026.
- [11] Peter Hase and Christopher Potts. Counterfactual simulation training for chain-of-thought faithfulness, 2026. URL <https://arxiv.org/abs/2602.20710>.
- [12] Usman Anwar, Tim Bakker, Dana Kianfar, Cristina Pinneri, and Christos Louizos. Analyzing and improving chain-of-thought monitorability through information theory. *arXiv preprint arXiv:2602.18297*, 2026.
- [13] Han Wang, Yifan Sun, Brian Ko, Mann Talati, Jiawen Gong, Zimeng Li, Naicheng Yu, Xucheng Yu, Wei Shen, Vedant Jolly, et al. Monitorbench: A comprehensive benchmark for chain-of-thought monitorability in large language models. *arXiv preprint arXiv:2603.28590*, 2026.
- [14] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- [15] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

- [16] Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, et al. Dapo: An open-source llm reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*, 2025.
- [17] Xu Shen, Song Wang, Zhen Tan, Laura Yao, Xinyu Zhao, Kaidi Xu, Xin Wang, and Tianlong Chen. Faithcot-bench: Benchmarking instance-level faithfulness of chain-of-thought reasoning. *arXiv preprint arXiv:2510.04040*, 2025.
- [18] Avni Mittal and Rauno Arike. C2-faith: Benchmarking llm judges for causal and coverage faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2603.05167*, 2026.
- [19] Zidi Xiong, Shan Chen, and Himabindu Lakkaraju. Monitorability as a free gift: How rlvr spontaneously aligns reasoning. *arXiv preprint arXiv:2602.03978*, 2026.
- [20] Chen Yueh-Han, Robert McCarthy, Bruce W Lee, He He, Ian Kivlichan, Bowen Baker, Micah Carroll, and Tomek Korbak. Reasoning models struggle to control their chains of thought. *arXiv preprint arXiv:2603.05706*, 2026.
- [21] Max Kaufmann, David Lindner, Roland S Zimmermann, et al. Aligned, orthogonal or in-conflict: When can we safely optimize chain-of-thought? *arXiv preprint arXiv:2603.30036*, 2026.
- [22] Miles Turpin, Andy Ardit, Marvin Li, Joe Benton, and Julian Michael. Teaching models to verbalize reward hacking in chain-of-thought reasoning. *arXiv preprint arXiv:2506.22777*, 2025.
- [23] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete problems in ai safety. *arXiv preprint arXiv:1606.06565*, 2016.
- [24] Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. *Advances in neural information processing systems*, 30, 2017.
- [25] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- [26] Joar Skalse, Nikolaus Howe, Dmitrii Krasheninnikov, and David Krueger. Defining and characterizing reward gaming. *Advances in Neural Information Processing Systems*, 35:9460–9471, 2022.
- [27] Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In *International Conference on Machine Learning*, pages 10835–10866. PMLR, 2023.
- [28] Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. Ai control: Improving safety despite intentional subversion. *arXiv preprint arXiv:2312.06942*, 2023.
- [29] Bowen Baker, Joost Huizinga, Leo Gao, Zehao Dou, Melody Y Guan, Aleksander Madry, Wojciech Zaremba, Jakub Pachocki, and David Farhi. Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. *arXiv preprint arXiv:2503.11926*, 2025.
- [30] Melody Y Guan, Miles Wang, Micah Carroll, Zehao Dou, Annie Y Wei, Marcus Williams, Benjamin Arnav, Joost Huizinga, Ian Kivlichan, Mia Glaese, et al. Monitoring monitorability. *arXiv preprint arXiv:2512.18311*, 2025.
- [31] David Saxton, Edward Grefenstette, Felix Hill, and Pushmeet Kohli. Analysing mathematical reasoning abilities of neural models. *arXiv preprint arXiv:1904.01557*, 2019.
- [32] Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, Thomas Hubert, Peter Choy, Cyprien de Masson d’Autume, Igor Babuschkin, Xinyun Chen, Po-Sen Huang, Johannes Welbl, Sven Gowal, Alexey Cherepanov, James Molloy, Daniel J. Mankowitz, Esme Sutherland Robson, Pushmeet Kohli, Nando de Freitas, Koray Kavukcuoglu, and Oriol Vinyals. Competition-level code generation with alphacode. *Science*, 378(6624):1092–1097, 2022. doi: 10.1126/science.abq1158. URL <https://www.science.org/doi/abs/10.1126/science.abq1158>.

- [33] Gemma Team Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, Tatiana Matejovicova, Alexandre Ram’e, Morgane Rivière, Louis Rouillard, Thomas Mesnard, Geoffrey Cideron, Jean-Bastien Grill, Sabela Ramos, Edouard Yvinec, Michelle Casbon, Etienne Pot, Ivo Penchev, Gael Liu, Francesco Visin, Kathleen Kenealy, Lucas Beyer, Xiaohai Zhai, Anton Tsitsulin, Róbert Istvan Busa-Fekete, Alex Feng, Noveen Sachdeva, Benjamin Coleman, Yi Gao, Basil Mustafa, Iain Barr, Emilio Parisotto, David Tian, Matan Eyal, Colin Cherry, Jan-Thorsten Peter, Danila Sinopalnikov, Surya Bhupatiraju, Rishabh Agarwal, Mehran Kazemi, Dan Malkin, Ravin Kumar, David Vilar, Idan Brusilovsky, Jiaming Luo, Andreas Steiner, Abe Friesen, Abhanshu Sharma, Abheesht Sharma, Adi Mayrav Gilady, Adrian Goedeckemeyer, Alaa Saade, Alexander Kolesnikov, Alexei Bendebury, Alvin Abdagic, Amit Vadi, Andr’as Gyorgy, André Susano Pinto, Anil Das, Ankur Bapna, Antoine Miech, Antoine Yang, Antonia Paterson, Ashish Shenoy, Ayan Chakrabarti, Bilal Piot, Boxi Wu, Bobak Shahriari, Bryce Pettrini, Charlie Chen, Charline Le Lan, Christopher A. Choquette-Choo, Cj Carey, Cormac Brick, Daniel Deutsch, Danielle Eisenbud, Dee Cattle, Derek Cheng, Dimitris Paparas, Divyashree Shivakumar Sreepathihalli, Doug Reid, Dustin Tran, Dustin Zelle, Eric Noland, Erwin Huizenga, Eugene Kharitonov, Frederick Liu, Gagik Amirkhanyan, Glenn Cameron, Hadi Hashemi, Hanna Klimczak-Plucińska, Harman Singh, Harsh Mehta, Harshal Tushar Lehri, Hussein Hazimeh, Ian Ballantyne, Idan Szpektor, Ivan Nardini, Jean Pouget-Abadie, Jetha Chan, Joe Stanton, J. Michael Wieting, Jonathan Lai, Jordi Orbay, Joe Fernandez, Joshua Newlan, Junsong Ji, Jyotinder Singh, Kat Black, Kathy Yu, Kevin Hui, Kiran Vodrahalli, Klaus Greff, Linhai Qiu, Marcella Valentine, Marina Coelho, Marvin Ritter, Matt Hoffman, Matthew Watson, Mayank Chaturvedi, Michael Moynihan, Min Ma, Nabila Babar, Natasha Noy, Nathan Byrd, Nick Roy, Nikola Momchev, Nilay Chauhan, Oskar Bunyan, Pankil Botarda, Paul Caron, Paul Kishan Rubenstein, Phil Culliton, Philipp Schmid, Pier Giuseppe Sessa, Ping mei Xu, Piotr Stańczyk, Pouya Dehghani Tafti, Rakesh Shivanna, Renjie Wu, Renke Pan, Reza Ardeshtir Rokni, Rob Willoughby, Rohith Vallu, Ryan Mullins, Sammy Jerome, Sara Smoot, Sertan Girgin, Shariq Iqbal, Shashir Reddy, Shruti Sheth, Siim Pöder, Sijal Bhatnagar, Sindhu Raghuram Panyam, Sivan Eiger, Susan Zhang, Tianqi Liu, Trevor Yacovone, Tyler Liechty, Uday Kalra, Utku Evci, Vedant Misra, Vincent Roseberry, Vladimir Feinberg, Vlad Kolesnikov, Woohyun Han, Woosuk Kwon, Xi Chen, Yinlam Chow, Yuvein Zhu, Zichuan Wei, Zoltan Egyed, Victor Cotruta, Minh Giang, Phoebe Kirk, Anand Rao, Jessica Lo, Erica Moreira, Luiz Gustavo Martins, Omar Sanseviero, Lucas Gonzalez, Zach Gleicher, Tris Warkentin, Vahab S. Mirrokni, Evan Senter, Eli Collins, Joelle Barral, Zoubin Ghahramani, Raia Hadsell, Yossi Matias, D. Sculley, Slav Petrov, Noah Fiedel, Noam Shazeer, Oriol Vinyals, Jeffrey Dean, Demis Hassabis, Koray Kavukcuoglu, Clément Farabet, Elena Buchatskaya, Jean-Baptiste Alayrac, Rohan Anil, Dmitry Lepikhin, Sebastian Borgeaud, Olivier Bachem, Armand Joulin, Alek Andreev, Cassidy Hardin, Robert Dadashi, and L’eonard Hussenot. Gemma 3 technical report. *ArXiv*, abs/2503.19786, 2025. URL <https://api.semanticscholar.org/CorpusID:277313563>.
- [34] Qwen Team. Qwen2.5: A party of foundation models, September 2024. URL <https://qwenlm.github.io/blog/qwen2.5/>.
- [35] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems*, 36:46595–46623, 2023.
- [36] Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. Hybridflow: A flexible and efficient rlhf framework. In *Proceedings of the Twentieth European Conference on Computer Systems*, pages 1279–1297, 2025.
- [37] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *arXiv preprint arXiv:1711.05101*, 2017.

Appendix

A Information flow based intervention methods Pipeline

Pipeline details. Figure A1 expands the intervention locations within the GRPO training loop. The top row shows the shared pipeline: the policy first generates K completions for each prompt, rewards are computed for each completion, group-normalized advantages are assigned, and the policy is updated with the GRPO objective. Our interventions do not modify rollout generation, reward computation, or advantage normalization. Thus, all methods are trained on the same sampled responses and reward signals.

The bottom rows show where each method changes the policy update. Vanilla GRPO uses full causal attention. Update Mask modifies the update-time forward pass by blocking answer-token attention to prompt tokens, directly removing the $A \rightarrow P$ shortcut during loss computation. Gradient Mask keeps the forward pass unchanged but stops the backward signal through the same $A \rightarrow P$ attention pathway. CoT-only Gradient leaves the forward computation numerically unchanged, but allows parameter gradients only from CoT positions, concentrating learning on the reasoning trace. FACT instead perturbs prompt hidden states during the update and trains the model under this perturbed representation, weakening prompt-side shortcut features. Together, these methods instantiate different ways of biasing the policy update toward the mediated path $P \rightarrow C \rightarrow A$.



Figure A1: **GRPO pipeline overview and intervention locations.** All methods share rollout generation, reward computation, and advantage normalization, and differ only in how they intervene during the policy update.

B Training details

All experiments use GRPO implemented in VERL [36], trained on a single node with 4 NVIDIA H200 GPUs. The `</think>` boundary is detected via substring matching, making delimiter identification tokenizer-agnostic.

Math task (DAPO-Math, Qwen2.5-7B-Instruct). We train Qwen2.5-7B-Instruct on DAPO-Math-17k using the AdamW [37] optimizer with a constant learning rate of 1×10^{-6} and a linear

warm-up over 10 gradient update steps. For rollout, the prompt batch size is 512 and we sample 8 responses per prompt with temperature 1.0 and top- p 1.0. The PPO mini-batch size is 32 with dynamic token-budget batching. The maximum prompt and response lengths are 1024 and 10,000 tokens, with no overlong penalty applied. We apply a KL regularization term with coefficient 0.001 and an entropy bonus with coefficient 0.001. For the Clip-Higher mechanism, we set $\varepsilon_{\text{low}} = 0.2$ and $\varepsilon_{\text{high}} = 0.28$. Training runs with gradient clipping at 1.0.

Code task (CodeContests reward hacking, Gemma3-4B-IT and Gemma3-12B-IT). We train gemma3-4b-it and gemma3-12b-it on the CodeContests reward-hacking dataset. For the 4B model we use a constant learning rate of $5e-7$. For rollout, the prompt batch size is 512 and we sample 4 responses per prompt with temperature 0.7 and top- p 0.9. The PPO mini-batch size is 128, yielding 16 gradient updates per rollout step. The maximum prompt and response lengths are 768 and 2,048 tokens. No KL regularization or entropy bonus is applied.

Math task (DeepMind Mathematics, Gemma3-1B). We train a gemma3-1b checkpoint on the full DeepMind Mathematics dataset using the AdamW optimizer with a cosine-decaying learning rate of 5×10^{-7} . For rollout, the prompt batch size is 512 and we sample 4 responses per prompt with temperature 0.7 and top- p 0.9. The PPO mini-batch size is 256, yielding 8 gradient updates per rollout step. The maximum prompt and response lengths are both 512 tokens. No KL regularization or entropy bonus is applied.

For FACT, we apply FGSM-based worst-case prompt perturbations at the embedding layer with $\epsilon = 0.05$ (4B code and 1B math) and $\epsilon = 0.2$ (12B code).

C Additional Experiments

C.1 A failure case for KL-based metrics




Figure A2: **KL direct effect becomes unstable in low-entropy regimes, while gradient direct effect remains more stable.**

KL-DE can be confounded when the answer distribution becomes nearly deterministic after RL convergence. In Figure A2, KL-DE on the buggy-code checkpoint from Section 6.3 shows a strong but non-monotonic dependence on full-context entropy H_{full} , suggesting that large KL values can partly reflect the instability of comparing peaked distributions rather than only shortcut reliance. Grad-DE varies more smoothly over the same range, so we use gradient-based metrics as the primary faithfulness measures in low-entropy regimes.

KL-based metrics still remain useful as complementary diagnostics. In particular, the reference-model variant can be computed with an open-source reference model q using the evaluated model’s generated prompt, CoT, and answer. Thus, although it is only a proxy for the evaluated model’s internal computation, it remains useful when the evaluated model’s weights or gradients are unavailable.

C.2 KL-based metrics

KL-based metrics for hinted math. Figure A3 reports KL-based completeness and necessity metrics on hinted arithmetic. For direct effect, both KL-DE and KL-DE_{ref} generally distinguish vanilla RL from the intervention methods, suggesting that the interventions reduce residual prompt-to-answer dependence. The reference-model variant provides a useful external check: it asks whether an open-source reference model can predict the answer from the CoT without relying directly on the prompt.

The necessity metrics are less reliable. In particular, $\text{KL-Nec}_{\text{ref}}$ reflects whether the reference model q changes its answer distribution when the CoT is removed, not whether the evaluated model itself relied on the CoT when producing the answer. If q can solve the task from the prompt alone, $\text{KL-Nec}_{\text{ref}}$ may remain small even when the evaluated model used its CoT; if q depends strongly on the generated CoT, it may be large even for post-hoc rationalizations. We therefore treat KL-based necessity, especially the reference-model variant, as an auxiliary diagnostic rather than direct evidence of model-internal CoT reliance.




Figure A3: KL-based faithfulness metrics on hinted arithmetic. We report direct effect and CoT necessity using both the evaluated model and an external reference model. Lower KL-DE indicates weaker direct prompt-to-answer dependence. Higher KL-Nec is intended to indicate stronger answer dependence on the CoT, but the reference-model necessity metric reflects the reference model’s use of the CoT rather than the evaluated model’s internal reliance.

C.3 LLM-judge results on code repair

To complement the rule-based verbalization metric in Section 6.3, we also evaluate CoT transparency using an LLM judge. The judge is asked whether the reasoning trace explicitly reveals the shortcut strategy used in the final code, such as hardcoding visible test cases or constructing a lookup table, rather than only describing a genuine bug fix.

Figure A4 shows two variants of this evaluation. Panel (A) reports the overall LLM-judge verbalization rate, while Panel (B) restricts evaluation to examples where the generated solution is classified as hacking. The trends are consistent with the rule-based results. Vanilla RL rapidly loses transparency: although it learns to exploit visible tests, its CoT almost never verbalizes the hacking strategy after training. Gradient Mask also becomes less transparent over training. In contrast, CoT Gradient maintains the highest and most stable verbalization rate, approaching near-complete disclosure on hacking examples. Update Mask and FACT also improve over vanilla, but are less stable than CoT Gradient.

These results confirm that the transparency gains are not an artifact of exact pattern matching. Even under a semantic LLM-judge evaluation, CoT Gradient most consistently makes reward-hacking behavior visible in the reasoning trace.




Figure A4: LLM-judge evaluation of CoT transparency on code repair. Panel (A) reports the overall verbalization rate judged by an LLM. Panel (B) reports the verbalization rate conditioned on hacking examples. Vanilla RL and Gradient Mask lose transparency over training, while CoT Gradient maintains the strongest and most stable verbalization of shortcut behavior.

C.4 Interventions scale to larger models on code repair

To verify that the conclusions of Section 6.3 are not specific to the 4B-scale results reported in the main text, we repeat the buggy-code repair experiment with GEMMA3-12B-IT under the same

training setup. The behavioral dynamics (Figure A5) and faithfulness metrics (Figure A6) reproduce the qualitative pattern of the 4B results.

As at the smaller scale, all methods rapidly saturate the visible test pass rate while the hidden test pass rate plateaus well below it, confirming that visible-test reward continues to be partially gamed by lookup-table shortcuts at 12B. Vanilla RL and Gradient Mask again almost never verbalize the lookup-table strategy, while CoT Gradient, Update Mask, and FACT substantially raise the CoT hack verbalization rate, with CoT Gradient remaining the most stable across training. The task-agnostic metrics agree: the same interventions that increase verbalization also yield lower $H(A | C)$, lower Grad-DE, and higher Grad-Nec relative to vanilla RL.

The main takeaway is that scaling from 4B to 12B parameters does not erase the faithfulness gap between vanilla RL and the structural interventions, nor does it change the relative ordering among interventions. CoT Gradient and FACT in particular continue to expose reward-hacking behavior in the trace and shift answer-relevant information onto the CoT, indicating that the structural mechanisms underlying the interventions transfer with scale rather than diminishing.




Figure A5: Behavioral dynamics of vanilla RL and faithfulness-oriented interventions on buggy-code fixing with GEMMA3-12B-IT. The four panels report visible test case pass rate, hidden test case pass rate, lookup-table hack rate, and CoT hack verbalization rate vs. training steps. Visible and hidden test pass rates measure reward optimization and generalization, respectively; lookup-table hack rate measures shortcut exploitation of visible tests; and CoT hack verbalization rate measures whether the shortcut strategy is exposed in the CoT. Higher CoT hack verbalization indicates better transparency of the model’s final code behavior.




Figure A6: Faithfulness metrics on buggy-code fixing with GEMMA3-12B-IT. The three panels report sufficiency, completeness, and necessity using $H(A | C)$, Grad-DE, and Grad-Nec across training steps. Better faithfulness corresponds to more sufficient CoT information, weaker direct prompt-to-answer reliance, and stronger answer dependence on the CoT.

C.5 Quantitative results

We provide aggregate results for the Gemma3-4B model used in the coding experiments in Section 6.3, complementing the qualitative examples in Figure A7. For each generated solution, we measure visible-test pass rate, hidden-test failure rate, lookup-table usage, and whether the lookup-table shortcut is verbalized in the model’s chain-of-thought.

The results show that all methods can obtain high visible-test reward by exploiting the provided assertions, while hidden-test failures reveal that many solutions rely on shortcut behavior rather than solving the underlying programming task. The main difference is CoT faithfulness: vanilla RL often uses a lookup-table answer while producing only generic debugging reasoning, so the shortcut is absent from the CoT. In contrast, FACT, CoT Gradient, and Update Mask more frequently expose the lookup-table construction inside the `<think>` trace.

These results suggest that the interventions do not necessarily eliminate reward hacking, but make the reward-relevant shortcut more visible in the reasoning trace. Among them, CoT Gradient provides a particularly strong trade-off, preserving visible-test performance while improving the faithfulness and auditability of the CoT.

D Limitations

Our study has several limitations. First, our proposed metrics are diagnostics rather than definitive causal tests of the model’s internal computation. Gradient-based measures provide a local sensitivity estimate of answer dependence on prompt and CoT tokens, while entropy- and KL-based metrics depend on the chosen model distribution, reference model, and attention-mask design. As discussed in Section 4, KL-based metrics can also be confounded in low-entropy regimes, where the answer distribution is nearly deterministic. Thus, these metrics should be interpreted as complementary evidence of CoT-mediated information flow rather than as complete measurements of faithfulness.

Second, our interventions improve the visibility of shortcut behavior but do not always eliminate the shortcut itself. In the code-repair setting, models can still learn lookup-table solutions that pass visible assertions while failing hidden tests; the key improvement is that several interventions make this behavior more likely to appear in the `<think>` trace. Similarly, in DAPO-Math, CoT Gradient makes wrong-hint reliance more verbalized when it occurs, but does not fully prevent the model from being influenced by unseen wrong hints. This distinction is important: our methods primarily improve CoT faithfulness and auditability, not necessarily robustness to every shortcut.

Third, our empirical scope is limited to the models, tasks, and training regimes studied here. We evaluate hinted arithmetic, reward-hackable code repair, and DAPO-Math, using models including Gemma3 and Qwen2.5. Although these settings cover different failure modes, they do not exhaust the space of reasoning tasks, model families, or RL objectives. Future work should test whether the same structural interventions scale to larger frontier models, longer-horizon agentic tasks, and settings where the shortcut is more subtle than explicit hint-following or lookup-table memorization.

Finally, some interventions introduce practical training costs. Update Mask and Gradient Mask require manipulating attention structure and are less compatible with efficient attention implementations, making them difficult to apply at larger scale. CoT Gradient and FACT are more scalable, but still add implementation complexity and may interact with model architecture, sequence length, and optimizer choices. A more complete understanding of these trade-offs is needed before applying structural faithfulness training as a general-purpose RL procedure.

E Broader impact

Our work aims to improve the auditability of chain-of-thought reasoning by making shortcut and reward-hacking behavior more visible in the reasoning trace. This may help model developers and auditors detect failures that would otherwise remain hidden behind plausible but unfaithful explanations. However, faithful CoT should not be treated as a complete safety guarantee: models may still rely on shortcuts, omit relevant information, or learn forms of obfuscation not covered by our evaluations. Our results should therefore be viewed as one component of a broader model-evaluation and oversight pipeline.

F Existing assets and licenses

We use publicly available datasets and models, including the DeepMind Mathematics Dataset, Code Contests, DAPO-Math-17K, Gemma3, and Qwen2.5. We cite the original creators of these assets in




Figure A7: **CoT faithfulness under reward hacking.** All four models receive the same buggy program and pass the visible assertions by hardcoding a lookup table. The key difference is whether the chain-of-thought faithfully exposes the strategy that determines the final answer. **FACT**, **CoT Gradient**, and **Update Mask** include the lookup-table construction inside their <think> traces, so the CoT reveals the shortcut used to obtain the visible-test reward. **Vanilla** also relies on a lookup-table answer, but its CoT contains only generic edge-case reasoning and never references the assertion inputs. This creates a faithfulness failure: the model’s final answer is produced by a shortcut that is absent from, and therefore not auditable through, the reasoning trace.

the main text and use them only for research purposes consistent with their released terms of use and licenses. No proprietary or private datasets are introduced in this work.

G LLM usage

We used LLM-based writing assistance for polishing, editing, and improving the clarity of the manuscript. LLMs were not used to generate experimental results.